

```
In [1]: 1 from sklearn.cluster import KMeans
        2 import pandas as pd
        3 import numpy as np
        4
```

## Load the file into a dataframe

```
In [2]: 1 FILE_NAME = 'housing.csv'
        2 datHousing = pd.read_csv(FILE_NAME)
        3
```

## Inspect the top rows

```
In [3]: 1 datHousing.head()
        2
```

```
Out[3]:
```

	CRIM	ZN	INDUS	RIVER	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT	MEDV
0	3.32105	0.0	19.58	Yes	0.871	5.403	100.0	1.3216	5	403	14.7	26.82	13.4
1	1.12658	0.0	19.58	Yes	0.871	5.012	88.0	1.6102	5	403	14.7	12.12	15.3
2	1.41385	0.0	19.58	Yes	0.871	6.129	96.0	1.7494	5	403	14.7	15.12	17.0
3	3.53501	0.0	19.58	Yes	0.871	6.152	82.6	1.7455	5	403	14.7	15.02	15.6
4	1.27346	0.0	19.58	Yes	0.605	6.250	92.6	1.7984	5	403	14.7	5.50	27.0

In [4]:

```

1 # Correctly get the unique values
2 r = datHousing['RIVER'].unique()
3 n = datHousing['NOX'].unique()
4 ra = datHousing['RAD'].unique()
5 t = datHousing['TAX'].unique()
6 p = datHousing['PRATIO'].unique()
7 z = datHousing['ZN'].unique()
8 c = datHousing['CRIM'].unique()
9
10 print(f"Unique values in RIVER: {r}")
11 print(f"Unique values in NOX: {n}")
12 print(f"Unique values in RAD: {ra}")
13 print(f"Unique values in TAX: {t}")
14 print(f"Unique values in PRATIO: {p}")
15 print(f"Unique values in ZN: {z}")
16 print(f"Unique values in CRIM: {c}")
17

```

Unique values in RIVER: ['Yes' 'No']

Unique values in NOX: [0.871 0.605 0.489 0.55 0.507 0.464 0.447 0.429 0.401 0.77

0.718 0.631 0.668 0.538 0.469 0.458 0.524 0.499 0.428 0.448  
0.439 0.41 0.403 0.411 0.453 0.4161 0.398 0.409 0.413 0.437  
0.426 0.449 0.445 0.52 0.547 0.581 0.624 0.51 0.488 0.422  
0.404 0.415 0.504 0.431 0.392 0.394 0.647 0.575 0.4 0.389  
0.385 0.405 0.433 0.472 0.544 0.493 0.46 0.4379 0.515 0.442  
0.518 0.484 0.429 0.435 0.671 0.7 0.693 0.659 0.597 0.679  
0.614 0.584 0.713 0.74 0.655 0.58 0.532 0.583 0.609 0.585  
0.573 ]

Unique values in RAD: [ 5 4 8 3 1 24 2 6 7]

Unique values in TAX: [403 277 276 307 223 254 216 198 666 296 242 222 311  
279 252 233 243 469

226 313 256 284 337 345 305 398 281 247 270 384 432 188 437 193 265 255  
329 402 348 224 300 330 315 244 264 285 241 293 245 289 358 304 287 430  
422 370 352 351 280 335 411 187 334 711 391 273]

Unique values in PRATIO: [14.7 18.6 16.4 17.4 17.6 14.9 13.6 20.2 15.3 17.8  
18.7 15.2 21. 19.2

10 2 17 0 16 0 21 1 17 0 15 1 10 7 16 1 10 0 10 10 5 10 0 10 0 0 0

## Dummy code the categoric variables

In [5]:

```

1 datHousing = pd.concat([datHousing, pd.get_dummies(datHousing['RIVER'], pr
2 datHousing.drop(['RIVER'], inplace=True, axis=1)
3
4

```

In [6]: 1 datHousing.head()

Out[6]:

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT	MEDV	RIVER_
0	3.32105	0.0	19.58	0.871	5.403	100.0	1.3216	5	403	14.7	26.82	13.4	
1	1.12658	0.0	19.58	0.871	5.012	88.0	1.6102	5	403	14.7	12.12	15.3	
2	1.41385	0.0	19.58	0.871	6.129	96.0	1.7494	5	403	14.7	15.12	17.0	
3	3.53501	0.0	19.58	0.871	6.152	82.6	1.7455	5	403	14.7	15.02	15.6	
4	1.27346	0.0	19.58	0.605	6.250	92.6	1.7984	5	403	14.7	5.50	27.0	

**3. For now, we are not going to break the data frame into training and testing partitions. Create a regression model on the entire data frame predicting MEDV using all of the other predictors. This is your baseline model. Report the R2 value and MSE for this model.**

### Create new dataset without target variable

In [7]: 1 # Create new dataset without the target variable which is in this case 'MEDV'  
2 datHousingSub = datHousing.drop(['MEDV'], axis=1)

In [8]: 1 datHousingSub.head()

Out[8]:

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT	RIVER_Yes
0	3.32105	0.0	19.58	0.871	5.403	100.0	1.3216	5	403	14.7	26.82	1
1	1.12658	0.0	19.58	0.871	5.012	88.0	1.6102	5	403	14.7	12.12	1
2	1.41385	0.0	19.58	0.871	6.129	96.0	1.7494	5	403	14.7	15.12	1
3	3.53501	0.0	19.58	0.871	6.152	82.6	1.7455	5	403	14.7	15.02	1
4	1.27346	0.0	19.58	0.605	6.250	92.6	1.7984	5	403	14.7	5.50	1

Create new dataset with target variable

```
In [9]: 1 target_name = list(['MEDV'])
        2 housing_target = datHousing[target_name]
        3
        4 print(datHousingSub.head(10))
        5 print(housing_target.head(10))
        6 print (len(datHousingSub))
        7 print (len(housing_target))
```

\	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT
0	3.32105	0.0	19.58	0.871	5.403	100.0	1.3216	5	403	14.7	26.82
1	1.12658	0.0	19.58	0.871	5.012	88.0	1.6102	5	403	14.7	12.12
2	1.41385	0.0	19.58	0.871	6.129	96.0	1.7494	5	403	14.7	15.12
3	3.53501	0.0	19.58	0.871	6.152	82.6	1.7455	5	403	14.7	15.02
4	1.27346	0.0	19.58	0.605	6.250	92.6	1.7984	5	403	14.7	5.50
5	1.83377	0.0	19.58	0.605	7.802	98.2	2.0407	5	403	14.7	1.92
6	1.51902	0.0	19.58	0.605	8.375	93.9	2.1620	5	403	14.7	3.32
7	0.13587	0.0	10.59	0.489	6.064	59.1	4.2392	4	277	18.6	14.66
8	0.43571	0.0	10.59	0.489	5.344	100.0	3.8750	4	277	18.6	23.09
9	0.17446	0.0	10.59	0.489	5.960	92.1	3.8771	4	277	18.6	17.27
RIVER_Yes											
0		1									
1		1									
2		1									
3		1									
4		1									
5		1									
6		1									
7		1									
8		1									
9		1									
MEDV											
0	13.4										
1	15.3										
2	17.0										
3	15.6										
4	27.0										
5	50.0										
6	50.0										
7	24.4										
8	20.0										
9	21.7										
506											
506											

```
In [10]: 1 from sklearn.model_selection import train_test_split
2
3 X_train, X_test, y_train, y_test = train_test_split(datHousingSub, housing
4 print(X_train.shape)
5 print(X_test.shape)
6
7 print(y_train.shape)
8 print(y_test.shape)
```

```
(354, 12)
(152, 12)
(354, 1)
(152, 1)
```

```
In [11]: 1 # import the LinearRegression class
2 from sklearn.linear_model import LinearRegression
3
4 regressor = LinearRegression(fit_intercept = True) # instantiate the Linea
5 regressor.fit(X_train, y_train) # train the model
6
7 print(f'r_sqr value: {regressor.score(X_train, y_train)}')
```

```
r_sqr value: 0.7231057672823237
```

```
In [12]: 1 y_train_pred = regressor.predict(X_train)
2 y_test_pred = regressor.predict(X_test)
```

```
In [13]: 1 #Report the R2 value and MSE for this model.
2 from sklearn.metrics import r2_score
3 from sklearn.metrics import mean_squared_error
4 from math import sqrt
5
6 print("Baseline Regression Model Results:")
7 print('RMSE train: %.3f, test: %.3f' % (
8     sqrt(mean_squared_error(y_train, y_train_pred)),
9     sqrt(mean_squared_error(y_test, y_test_pred))))
10 print('R^2 train: %.3f, test: %.3f' % (
11     r2_score(y_train, y_train_pred),
12     r2_score(y_test, y_test_pred)))
```

```
Baseline Regression Model Results:
RMSE train: 4.730, test: 4.843
R^2 train: 0.723, test: 0.746
```

```
In [ ]: 1
```

## Fit 2 cluster model

**4. Create a data frame called `datHousingSub` (ALREADY DONE) for clustering by generating a subset with all columns except `MEDV`. We want to exclude `MEDV` from the clustering since this is the value that we will try to predict later from the clusters. Do a summary on the subset to verify.**

```
In [14]: 1 # Using scikit-learn to perform K-Means clustering
        2
        3 # Specify the number of clusters (2) and fit the data dat_rider_feats
        4 kmeans = KMeans(n_clusters=2, random_state=42).fit(datHousingSub)
        5
```

C:\anaconda3\lib\site-packages\sklearn\cluster\\_kmeans.py:1416: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning  
 super().\_check\_params\_vs\_input(X, default\_n\_init=10)

```
In [15]: 1 # Set global NumPy print options for formatting
        2 np.set_printoptions(suppress=True)
        3
        4 print(datHousingSub.columns)
        5 centroids = (kmeans.cluster_centers_)
        6 print(centroids)
```

```
Index(['CRIM', 'ZN', 'INDUS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
       'PRATIO', 'LSTAT', 'RIVER_Yes'],
      dtype='object')
[[ 0.38877444 15.58265583  8.42089431  0.51184743  6.38800542
  60.63224932  4.44127154  4.45528455 311.92682927 17.80921409
  10.41745257  0.07317073]
 [12.29916168  0.          18.45182482  0.67010219  6.00621168
  89.96788321  2.05447007 23.27007299 667.64233577 20.19635036
 18.67452555  0.05839416]]
```

**5. Using the KMeans algorithm, create a model with 2 clusters on the datHousingSub data frame. Report the centroid values for each cluster and the sizes of each cluster. Using the characteristics that are especially divergent between the two clusters, what would you name these clusters?**

*To me this looks like 0 is group low crime, and higher Zonning number (less pop per acre), with Lower Industrial activity, and lower age, I would assess this to represent a Suburban center which tend to be newer areas, with lower crime rates than urban cities. We can see the opposite of this is true for group 1, as it is Older and have higher crime rate, 0 for zoning and industrial activity is rather high in comparison.*

**0 = Suburban and 1 = Urban**



In [22]:

```
1 k_labels = (k_means.labels_)
2 print(k_labels)
```

[illegible]

In [23]:

```
1 # Calculate silhouette_score
2 from sklearn.metrics import silhouette_score
3
4 print(silhouette_score(datHousingSub, k_means.labels_))
```

0.6217529916045139

***Do you think that adding another cluster helps to partition the data? Why or why not?***

No not really, although I would call the 3 groups in order 0 = Urban, 1 = Rural, and 2 = Suburban based on the variables mentioned before, where the groups seem to be split by zoning and industrial activities, which give indication in pop per acre and how many business are close by, this is also supported by the DIS variable for distance to work. However, based on the silhouette\_score I would say the 2 cluster test is sufficient, this is purely based on the fact between what I classified as rural and urban would be very small, and mostly likely closer to urban and outer urban area as the variable seem to be slightly close and would give better explanation power if they were combined or in a cluster of 2. The lower silhouette score indicates that the clusters are closer to overlapping than the previous score.

In [25]:

```
1 kmeans
```

```
Out[25]: KMeans(n_clusters=2, random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

On GitHub, the HTML representation is unable to render, please try loading this page with [nbviewer.org](https://nbviewer.org).



In [26]: 1 centroids

Out[26]: array([[ 0.38877444, 15.58265583, 8.42089431, 0.51184743,  
6.38800542, 60.63224932, 4.44127154, 4.45528455,  
311.92682927, 17.80921409, 10.41745257, 0.07317073],  
[ 12.29916168, 0. , 18.45182482, 0.67010219,  
6.00621168, 89.96788321, 2.05447007, 23.27007299,  
667.64233577, 20.19635036, 18.67452555, 0.05839416]])

**7. We will use the 2-cluster model going forward. Merge the cluster ids from this 2-cluster model into the datHousing data frame using the column name Cluster to store this cluster id. Look at the first few rows of data to verify.**

In [36]: 1 # Merge the existing 'labels' variable directly into the DataFrame  
2 datHousing['Cluster'] = labels  
3  
4 # Verify the first few rows  
5 print(datHousing.head())  
6 print(datHousing['Cluster'].unique())

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT
0	3.32105	0.0	19.58	0.871	5.403	100.0	1.3216	5	403	14.7	26.82
1	1.12658	0.0	19.58	0.871	5.012	88.0	1.6102	5	403	14.7	12.12
2	1.41385	0.0	19.58	0.871	6.129	96.0	1.7494	5	403	14.7	15.12
3	3.53501	0.0	19.58	0.871	6.152	82.6	1.7455	5	403	14.7	15.02
4	1.27346	0.0	19.58	0.605	6.250	92.6	1.7984	5	403	14.7	5.50

  

	MEDV	RIVER_Yes	Cluster
0	13.4	1	0
1	15.3	1	0
2	17.0	1	0
3	15.6	1	0
4	27.0	1	0

[0 1]

**8. Create a new data frame called datHousingC1 which contains all of the rows from datHousing in cluster 1. Look at the first few rows of data to verify. Check the Cluster column to make sure that it only stores the value of 1**

```
In [51]: 1 datHousingC1 = datHousing.loc[datHousing['Cluster'] == 0]
2 # Verify the first few rows
3 print(datHousingC1.head())
4 print(datHousingC1['Cluster'].unique())
5 print(len(datHousingC1))
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT
\											
0	3.32105	0.0	19.58	0.871	5.403	100.0	1.3216	5	403	14.7	26.82
1	1.12658	0.0	19.58	0.871	5.012	88.0	1.6102	5	403	14.7	12.12
2	1.41385	0.0	19.58	0.871	6.129	96.0	1.7494	5	403	14.7	15.12
3	3.53501	0.0	19.58	0.871	6.152	82.6	1.7455	5	403	14.7	15.02
4	1.27346	0.0	19.58	0.605	6.250	92.6	1.7984	5	403	14.7	5.50

	MEDV	RIVER_Yes	Cluster
0	13.4	1	0
1	15.3	1	0
2	17.0	1	0
3	15.6	1	0
4	27.0	1	0

[0]  
369

**9. Create a new data frame called datHousingC2 which contains all of the rows from datHousing in cluster 2. Use summary to verify the results. Check the Cluster column to make sure that it only stores the value of 2**

```
In [50]: 1 datHousingC2 = datHousing.loc[datHousing['Cluster'] == 1]
2 # Verify the first few rows
3 print(datHousingC2.head())
4 print(datHousingC2['Cluster'].unique())
5 print(len(datHousingC2))
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT
\											
27	8.98296	0.0	18.1	0.770	6.212	97.4	2.1222	24	666	20.2	17.60
28	3.84970	0.0	18.1	0.770	6.395	91.0	2.5052	24	666	20.2	13.27
29	5.20177	0.0	18.1	0.770	6.127	83.4	2.7227	24	666	20.2	11.48
30	4.22239	0.0	18.1	0.770	5.803	89.0	1.9047	24	666	20.2	14.64
31	3.47428	0.0	18.1	0.718	8.780	82.9	1.9047	24	666	20.2	5.29

	MEDV	RIVER_Yes	Cluster
27	17.8	1	1
28	21.7	1	1
29	22.7	1	1
30	16.8	1	1
31	21.9	1	1

[1]  
137

**10. Create a regression model predicting MEDV and the same predictors as the baseline model in step 3, using the data frame from cluster 1. What are the R2 value and MSE? Is this higher or lower than the baseline model?**

```
In [52]: 1 target_name = list(['MEDV'])
2 housing_target1 = datHousingC1[target_name]
3 datHousingSub1 = datHousingC1.drop(['MEDV', 'Cluster'], axis=1)
4
5 print(datHousingSub1.head(10))
6 print(housing_target1.head(10))
7 print (len(datHousingSub1))
8 print (len(housing_target1))
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT
\											
0	3.32105	0.0	19.58	0.871	5.403	100.0	1.3216	5	403	14.7	26.82
1	1.12658	0.0	19.58	0.871	5.012	88.0	1.6102	5	403	14.7	12.12
2	1.41385	0.0	19.58	0.871	6.129	96.0	1.7494	5	403	14.7	15.12
3	3.53501	0.0	19.58	0.871	6.152	82.6	1.7455	5	403	14.7	15.02
4	1.27346	0.0	19.58	0.605	6.250	92.6	1.7984	5	403	14.7	5.50
5	1.83377	0.0	19.58	0.605	7.802	98.2	2.0407	5	403	14.7	1.92
6	1.51902	0.0	19.58	0.605	8.375	93.9	2.1620	5	403	14.7	3.32
7	0.13587	0.0	10.59	0.489	6.064	59.1	4.2392	4	277	18.6	14.66
8	0.43571	0.0	10.59	0.489	5.344	100.0	3.8750	4	277	18.6	23.09
9	0.17446	0.0	10.59	0.489	5.960	92.1	3.8771	4	277	18.6	17.27

RIVER_Yes	
0	1
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	1
9	1

MEDV	
0	13.4
1	15.3
2	17.0
3	15.6
4	27.0
5	50.0
6	50.0
7	24.4
8	20.0
9	21.7
369	
369	

```
In [58]: 1 X_train, X_test, y_train, y_test = train_test_split(datHousingSub1, housin
2 print('shape of training data independent vars (x)')
3 print(X_train.shape)
4 print('shape of test data independent vars (x)')
5 print(X_test.shape)
6 print('shape of train data dependent vars (y)')
7 print(y_train.shape)
8 print('shape of test data dependent vars (y)')
9 print(y_test.shape)
10 regressor = LinearRegression(fit_intercept = True) # instantiate the Linea
11 regressor.fit(X_train, y_train) # train the model
12 print(f'r_sqr value: {regressor.score(X_train, y_train)}')
13 y_train_pred = regressor.predict(X_train)
14 y_test_pred = regressor.predict(X_test)
15 print("Cluster1 Regression Model Results:")
16 print('RMSE train: %.3f, test: %.3f' % (
17     sqrt(mean_squared_error(y_train, y_train_pred)),
18     sqrt(mean_squared_error(y_test, y_test_pred))))
19 print('R^2 train: %.3f, test: %.3f' % (
20     r2_score(y_train, y_train_pred),
21     r2_score(y_test, y_test_pred)))
```

shape of training data independent vars (x)

(258, 12)

shape of test data independent vars (x)

(111, 12)

shape of train data dependent vars (y)

(258, 1)

shape of test data dependent vars (y)

(111, 1)

r\_sqr value: 0.8477216916255111

Cluster1 Regression Model Results:

RMSE train: 3.128, test: 3.121

R^2 train: 0.848, test: 0.877

### **Baseline Regression Model Results:**

**RMSE train: 4.730, test: 4.843**

**R^2 train: 0.723, test: 0.746**

What are the R2 value and MSE? Is this higher or lower than the baseline model? These are higher than base line model as we can see this model performs better at accurately predicting the data relevant to cluster 1 which in this case, would be the data we had a larger amount of our "suburban area = 0" (maybe explaining the higher predictive power)

In [ ]:

1

**11. Create a regression model predicting MEDV and the same predictors as the baseline model in step 3, using the data frame from cluster 2. What are the R2 value and MSE? Is this higher or lower than the baseline model?**

```
In [60]: 1 target_name = list(['MEDV'])
2 housing_target2 = datHousingC2[target_name]
3 datHousingSub2 = datHousingC2.drop(['MEDV', 'Cluster'], axis=1)
4
5 print(datHousingSub2.head(10))
6 print(housing_target2.head(10))
7 print (len(datHousingSub2))
8 print (len(housing_target2))
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PRATIO	LSTAT
\											
27	8.98296	0.0	18.1	0.770	6.212	97.4	2.1222	24	666	20.2	17.60
28	3.84970	0.0	18.1	0.770	6.395	91.0	2.5052	24	666	20.2	13.27
29	5.20177	0.0	18.1	0.770	6.127	83.4	2.7227	24	666	20.2	11.48
30	4.22239	0.0	18.1	0.770	5.803	89.0	1.9047	24	666	20.2	14.64
31	3.47428	0.0	18.1	0.718	8.780	82.9	1.9047	24	666	20.2	5.29
32	5.66998	0.0	18.1	0.631	6.683	96.8	1.3567	24	666	20.2	3.73
33	6.53876	0.0	18.1	0.631	7.016	97.5	1.2024	24	666	20.2	2.96
34	8.26725	0.0	18.1	0.668	5.875	89.6	1.1296	24	666	20.2	8.88
364	4.26131	0.0	18.1	0.770	6.112	81.3	2.5091	24	666	20.2	12.67
365	4.54192	0.0	18.1	0.770	6.398	88.0	2.5182	24	666	20.2	7.79
RIVER_Yes											
27		1									
28		1									
29		1									
30		1									
31		1									
32		1									
33		1									
34		1									
364		0									
365		0									
MEDV											
27	17.8										
28	21.7										
29	22.7										
30	16.8										
31	21.9										
32	50.0										
33	50.0										
34	50.0										
364	22.6										
365	25.0										
137											
137											

```
In [61]: 1 X_train, X_test, y_train, y_test = train_test_split(datHousingSub2, housin
2 print('shape of training data independent vars (x)')
3 print(X_train.shape)
4 print('shape of test data independent vars (x)')
5 print(X_test.shape)
6 print('shape of train data dependent vars (y)')
7 print(y_train.shape)
8 print('shape of test data dependent vars (y)')
9 print(y_test.shape)
10 regressor = LinearRegression(fit_intercept = True) # instantiate the Linea
11 regressor.fit(X_train, y_train) # train the model
12 print(f'r_sqr value: {regressor.score(X_train, y_train)}')
13 y_train_pred = regressor.predict(X_train)
14 y_test_pred = regressor.predict(X_test)
15 print("Cluster1 Regression Model Results:")
16 print('RMSE train: %.3f, test: %.3f' % (
17     sqrt(mean_squared_error(y_train, y_train_pred)),
18     sqrt(mean_squared_error(y_test, y_test_pred))))
19 print('R^2 train: %.3f, test: %.3f' % (
20     r2_score(y_train, y_train_pred),
21     r2_score(y_test, y_test_pred)))
```

shape of training data independent vars (x)

(95, 12)

shape of test data independent vars (x)

(42, 12)

shape of train data dependent vars (y)

(95, 1)

shape of test data dependent vars (y)

(42, 1)

r\_sqr value: 0.6735080000289209

Cluster1 Regression Model Results:

RMSE train: 4.709, test: 5.020

R^2 train: 0.674, test: 0.677

### **Baseline Regression Model Results:**

**RMSE train: 4.730, test: 4.843**

**R^2 train: 0.723, test: 0.746**

What are the R2 value and MSE? Is this higher or lower than the baseline model? These are lower than base line model as we can see this model performs worse than the baseline at accurately predicting the data relevant to cluster 2 which in this case, would be the data we had a less amount of our "urban area = 1".

### **12. Summarize your findings. Do you think that clustering might improve your ability to predict the MEDV value? If so, under what contexts or constraints?**

Clustering can improve the predictive power of a model as we can see in this example which may be due to the amount of data points vs 0 and 1 cluster giving better predictive power to the underlying factors that influence the target variable (MEDV) within different clusters. In this case:

For suburban areas, the clustering improved the ability to predict MEDV.

For urban areas, where the data maybe be more complex or less of, clustering did not improve the model's predictive performance.

In [ ]:

1